

Rishabh Chavan

Data Science Intern @Bharat Intern

Task 1: Stock Price Prediction using LSTM Model

Importing Libraries

```
In [1]: import yfinance as yf
import pandas as pd
import numpy as np
import tensorflow as tf
import pandas_datareader as web
import os
from keras.models import Sequential
from keras.layers import LSTM, Dense
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
plt.style.use('f5vethurlyright')
pd.options.mode.chained_assignment = None

In [2]: df = pd.read_csv('C:\Users\LENOVO\OneDrive\Documents\TSLA.csv')

In [3]: os.getcwd()

Out[3]: 'C:\Users\LENOVO'

In [4]: os.chdir('C:\Users\LENOVO\OneDrive\Documents')

In [5]: df = pd.read_csv('TSLA.csv')


In [6]: df

Out[6]:
```

	Date	Open	High	Low	Close	Adj Close	Volume
0	2018-05-15	19.000668	19.130667	18.700001	18.945333	18.945333	142768000
1	2018-05-16	18.922001	19.254000	18.770666	19.098667	19.098667	85110000
2	2018-05-17	19.059999	19.279333	18.931334	18.969334	18.969334	66390000
3	2018-05-18	18.976667	18.976667	18.266666	18.454666	18.454666	108779500
4	2018-05-21	18.755333	19.432667	18.753332	18.966000	18.966000	137739000
...
1253	2023-05-08	173.720001	173.800003	169.190002	171.789993	171.789993	112249400
1254	2023-05-09	168.949997	169.820007	166.559998	169.149994	169.149994	88965000
1255	2023-05-10	172.550003	174.429993	166.679993	168.539993	168.539993	119840700
1256	2023-05-11	168.699997	173.570007	166.789993	172.080002	172.080002	103889900
1257	2023-05-12	176.070007	177.380005	167.229996	167.979996	167.979996	157577100

1258 rows x 7 columns

Visualize the Closing Price



Create a new dataframe with only Close Column

```
In [8]: data = df.filter(['Close'])
```

Convert the dataframe to numpy array

```
In [9]: dataset = data.values
```

Train 80 percent of number of rows

```
In [10]: training_data_len = int(np.ceil( len(dataset) *0.8 ))
training_data_len

Out[10]: 1007
```

Scale the data

```
In [11]: scaler = StandardScaler()
scaled_data = scaler.fit_transform(dataset)
print("MEAN of processed data: ",scaled_data.mean())
print("Standard Deviation of processed data: ",scaled_data.std())

MEAN of processed data:  0.0
Standard Deviation of processed data:  1.0
```

Split the data into x_train and y_train

```
In [12]: train_data = scaled_data[0:int(training_data_len), :]
```

```
In [13]: x_train = []
y_train = []
for i in range(60, len(train_data)):
    x_train.append(train_data[i-60:1, 0])
    y_train.append(train_data[i, 0])
```

Convert the x_train and y_train to numpy arrays

```
In [14]: x_train, y_train = np.array(x_train), np.array(y_train)
```

Reshape the data - x_train

```
In [15]: x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))
print(x_train.shape)

(947, 60, 1)
```

Build LSTM Model

```
In [16]: model = Sequential()
model.add(LSTM(100, return_sequences=True, input_shape= (x_train.shape[1], 1)))
model.add(LSTM(100, return_sequences=False))
model.add(Dense(50))
model.add(Dense(25))
model.add(Dense(1))
```

Compile the model

```
In [17]: model.compile(optimizer='adam', loss='mean_squared_error')
model.summary()

Model: "sequential"

Layer (type)                 Output Shape         Param #
=====
lstm (LSTM)                  (None, 60, 100)     48800
lstm_1 (LSTM)                (None, 100)         88400
dense (Dense)                (None, 50)          5050
dense_1 (Dense)              (None, 25)          1275
dense_2 (Dense)              (None, 1)            26
=====
Total params: 127,551
Trainable params: 127,551
Non-trainable params: 0
```

Train the model

```
In [18]: history = model.fit(x_train, y_train, epochs=10, batch_size=32)

Epoch 1/10
30/30 [=====] - 25s 179ms/step - loss: 0.1214
Epoch 2/10
30/30 [=====] - 5s 160ms/step - loss: 0.0163
Epoch 3/10
30/30 [=====] - 5s 154ms/step - loss: 0.0154
Epoch 4/10
30/30 [=====] - 4s 146ms/step - loss: 0.0154
Epoch 5/10
30/30 [=====] - 5s 151ms/step - loss: 0.0138
Epoch 6/10
30/30 [=====] - 4s 147ms/step - loss: 0.0113
Epoch 7/10
30/30 [=====] - 5s 166ms/step - loss: 0.0107
Epoch 8/10
30/30 [=====] - 5s 160ms/step - loss: 0.0120
Epoch 9/10
30/30 [=====] - 4s 138ms/step - loss: 0.0091
Epoch 10/10
30/30 [=====] - 5s 153ms/step - loss: 0.0085

In [19]: plt.plot(history.history['loss'])

Out[19]: [C:\Users\LENOVO\OneDrive\Documents\TSLA.csv]
```



Create the data - x_test and y_test

```
In [20]: test_data = scaled_data[training_data_len - 60: , :]
```

```
In [21]: x_test = []
y_test = dataset[training_data_len:, :]
for i in range(60, len(test_data)):
    x_test.append(test_data[i-60:1, 0])
```

Convert the data to a numpy array

```
In [22]: x_test = np.array(x_test)
```

Reshape the data - x_test

```
In [23]: x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1 ))
```

Evaluate the model

```
In [24]: train_loss = model.evaluate(x_train, y_train, verbose=0)
train_loss = model.evaluate(x_test, y_test, verbose=0)
print(f'Train Loss: {train_loss:.4f}')
print(f'Test Loss: {train_loss:.4f}')

Train Loss: 47550.5977
Test Loss: 47550.5977
```

Price Value Prediction

```
In [25]: predictions = model.predict(x_test)
predictions = scaler.inverse_transform(predictions)

8/8 [=====] - 4s 58ms/step
```

Calculate RMSE

```
In [26]: rmse = np.sqrt(np.mean((predictions - y_test) ** 2))
rmse

Out[26]: 12.016287234461927
```

Plot the data

```
In [27]: train = data[:training_data_len]
valid = data[training_data_len:]
valid['Predictions'] = predictions

In [28]: plt.figure(figsize=(16,6))
plt.title('Prediction of Closing Price')
plt.xlabel('Time')
plt.ylabel('Close Price ($)')
plt.plot(train['Close'])
plt.plot(valid[['Close', 'Predictions']])
plt.legend(['Train', 'Test', 'Predictions'], loc='lower right')
plt.show()
```



Show the valid and predicted prices

```
In [29]: valid

Out[29]:
```

	Close	Predictions
1007	256.529999	254.450226
1008	241.456665	249.151672
1009	253.869995	243.889145
1010	236.603333	242.330399
1011	236.473328	239.477356
...
1253	171.789993	162.366913
1254	169.149994	164.703659
1255	168.539993	166.800720
1256	172.080002	168.441315

1257 167.979996 170.140335
251 rows x 2 columns